

# Machine Learning in Production Safety

# Mitigating more mistakes...

## Fundamentals of Engineering AI-Enabled Systems

**Holistic system view:** AI and non-AI components, pipelines, stakeholders, environment interactions, feedback loops

### Requirements:

System and model goals  
User requirements  
Environment assumptions  
Quality beyond accuracy  
Measurement  
Risk analysis  
Planning for mistakes

### Architecture + design:

Modeling tradeoffs  
Deployment architecture  
Data science pipelines  
Telemetry, monitoring  
Anticipating evolution  
Big data processing  
Human-AI design

### Quality assurance:

Model testing  
Data quality  
QA automation  
Testing in production  
Infrastructure quality  
Debugging

### Operations:

Continuous deployment  
Contin. experimentation  
Configuration mgmt.  
Monitoring  
Versioning  
Big data  
DevOps, MLOps

**Teams and process:** Data science vs software eng. workflows, interdisciplinary teams, collaboration points, technical debt

## Responsible AI Engineering

Provenance,  
versioning,  
reproducibility

Safety

Security and  
privacy

Fairness

Interpretability  
and explainability

Transparency  
and trust

Ethics, governance, regulation, compliance, organizational culture

# Reading

S. Mohseni et al., Practical Solutions for Machine Learning Safety in Autonomous Vehicles. SafeAI Workshop@AAAI (2020).

# Learning Goals

- Understand safety concerns in traditional and AI-enabled systems
- Understand the importance of ML robustness for safety
- Apply hazard analysis to identify risks and requirements and understand their limitations
- Discuss ways to design systems to be safe against potential failures
- Suggest safety assurance strategies for a specific project
- Describe the typical processes for safety evaluations and their limitations

# System Safety

# Defining Safety

Prevention of a system failure or malfunction that results in:

- Death or serious injury to people
- Loss or severe damage to equipment/property
- Harm to the environment or society

Safety is a system concept

- Can't talk about software/ML being "safe"/"unsafe" on its own
- Safety is defined in terms of its effect on the **environment**

# Safety != Reliability

Reliability = absence of defects, mean time between failure

Safety = prevents accidents, harms

Can build safe systems from unreliable components (e.g. redundancy, safeguards)

System may be unsafe despite reliable components (e.g. stronger gas tank causes more severe damage in incident)

Accuracy is usually about reliability!

# Safety of AI-Enabled Systems



Dr. Emily Slackerman Ackerman  
@EmilyEAckerman · [Follow](#)



i (in a wheelchair) was just trapped \*on\* forbes ave by one of these robots, only days after their independent roll out. i can tell that as long as they continue to operate, they are going to be a major accessibility and safety issue. [thread]



[pittnews.com](http://pittnews.com)  
Everything we know about the Starship food delivery robots  
The white, 2-foot tall battery-powered delivery robots will be...

7:27 PM · Oct 21, 2019



[Read the full conversation on Twitter](#)



4K

Reply

Copy link

## Speaker notes

Systems can be unsafe in unexpected ways



# Safety of AI-Enabled Systems

The screenshot shows a tweet from user **skoops** (@skoops) with three facepalming emojis. The tweet reads: "The [@netatmo](#) servers are down and twitter is already full of freezing people not able to control their heating :D (via [protected]) / cc [@internetofshit](#)". Below the tweet are two replies:

**Kieran** (@DivemasterK) - "Are your servers down ? I can connect to my app to turn on heating !! 2.11.18, 21:02 from Wicklow, Ireland"

**Kiran vadgama** (@kiran\_vadgama) - "@netatmo hi my manual ovrride on my thermostat is not working and when i try using the app it comes up with an error with servers down. Can i overide at boiler end?"

**Andy Mc** (@ITakeSugar) - "Replies to @levisleedaniel and @netatmo Is there a way to control the boiler even if the servers are down, it's freezing at the moment 22.11.18, 20:38"

**James Brown** (@jamesbrun) - "Replies to @tyrestighe @levisleedaniel @netatmo same issue. Can't control heating via app cannot login to [netatmo.com](#) to control from there. What is the status @netatmo ?"

At the bottom of the tweet, there are engagement metrics: 8:15 PM · Nov 22, 2018, 2K likes, 1 reply, and a link to "Read 67 replies".

## Speaker notes

Systems can be unsafe in unexpected ways



# Safety is a broad concept

Not just physical harms/injuries to people

Includes harm to mental health

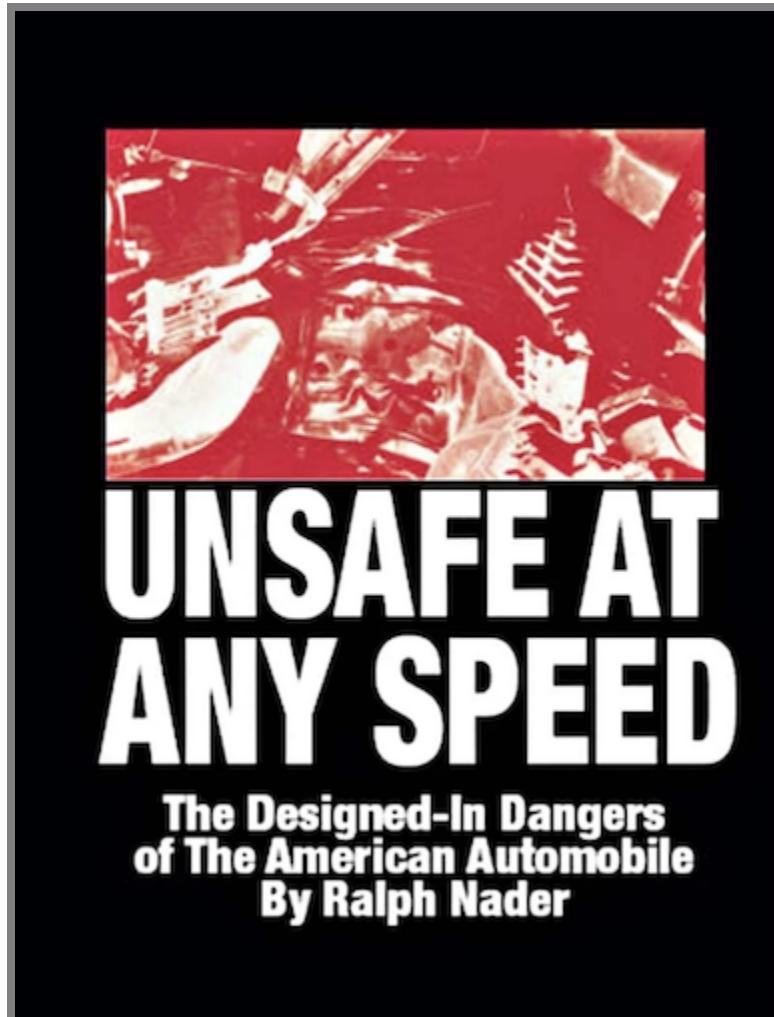
Includes polluting the environment, including noise pollution

Includes harm to society, e.g. poverty, polarization

# Case Study: Self-Driving Car



# How did traditional vehicles become safer?



National Traffic & Motor Safety Act (1966):

- Mandatory design changes (head rests, shatter-resistant windshields, safety belts)
- Road improvements (center lines, reflectors, guardrails)
- Significant reduction (13-46%) in traffic fatalities

# Autonomous Vehicles: What's different?

## Ford Taps the Brakes on the Arrival of Self-Driving Cars

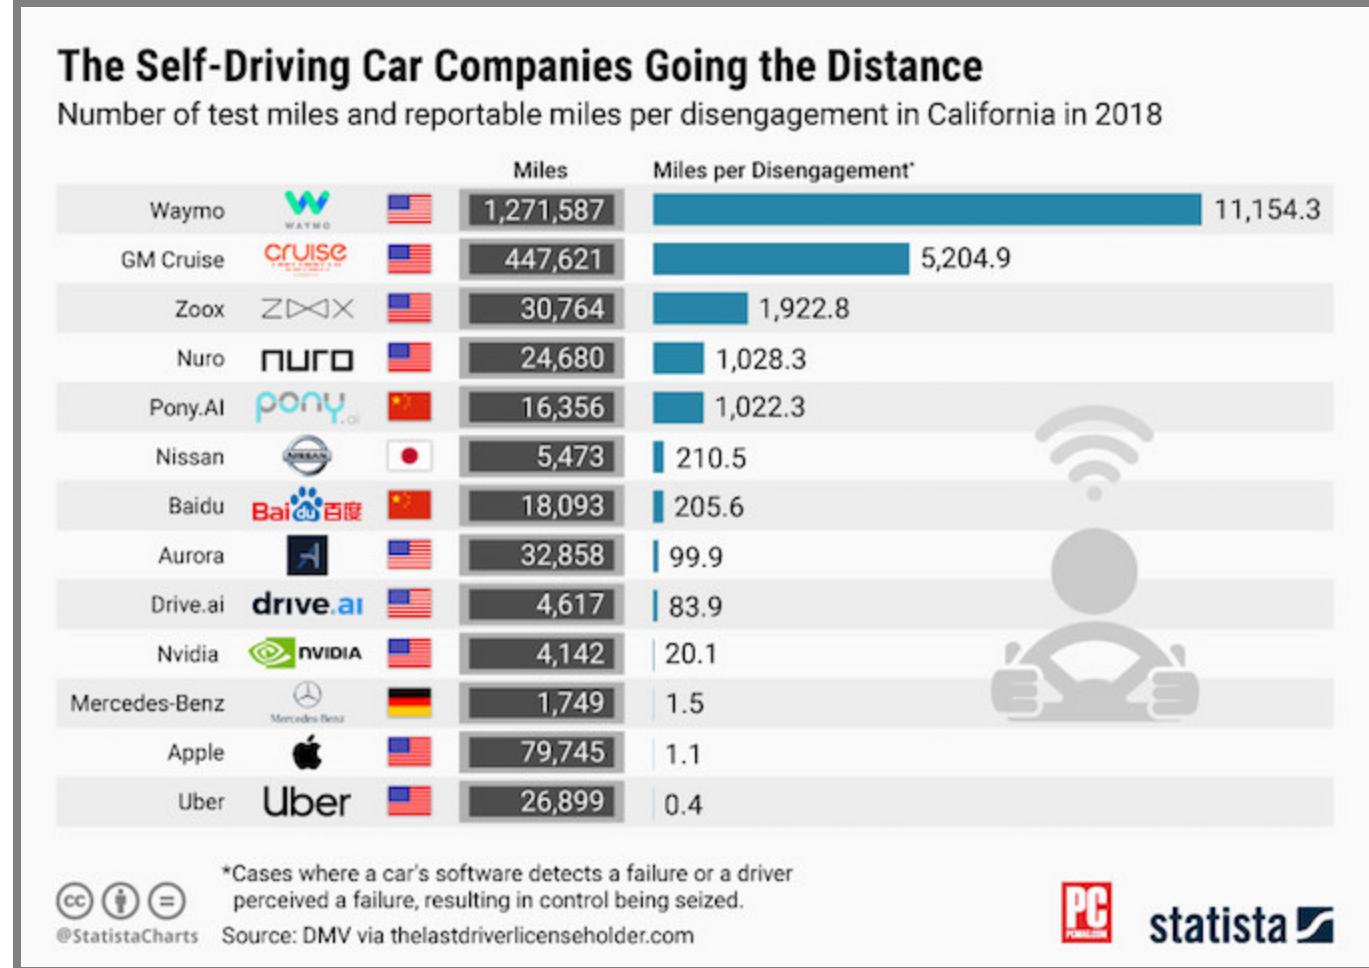
HYPE CYCLE —

The hype around driverless cars came crashing down in 2018

Top Toyota expert throws cold water on the driverless car hype

- In traditional vehicles, humans ultimately responsible for safety
  - Built-in safety features (lane keeping, emergency braking)
  - i.e., safety = human control + safety mechanisms
- Use of AI in autonomous vehicles: Perception, control, routing, etc.,
  - Inductive training: No explicit requirements or design insights
  - Can ML achieve safe design solely through lots of data?

# Demonstrating Safety



≡ Q. More miles tested => safer?

# Challenge: Edge/Unknown Cases



- Gaps in training data; ML unlikely to cover all unknown cases
- Is this a unique problem for AI? What about humans?

# Safety Engineering

Safety Engineering: An engineering discipline which assures that engineered systems provide acceptable levels of safety.

Typical safety engineering process:

- Identify relevant hazards & safety requirements
- Identify potential root causes for hazards
- For each hazard, develop a mitigation strategy
- Provide evidence that mitigations are properly implemented

# Improving Safety of ML-Enabled Systems

Anticipate problems (hazard analysis, FTA, FMEA, HAZOP, ...)

Anticipate the existence of unanticipated problems

Plan for mistakes, design mitigations (recall earlier lecture!)

- Human in the loop
- Undoable actions, failsoft
- Guardrails
- Mistaked detection
- Redundancy, ...

# Demonstrating and Documenting Safety

# Demonstrating Safety

Two main strategies:

**1. Evidence of safe behavior in the field**

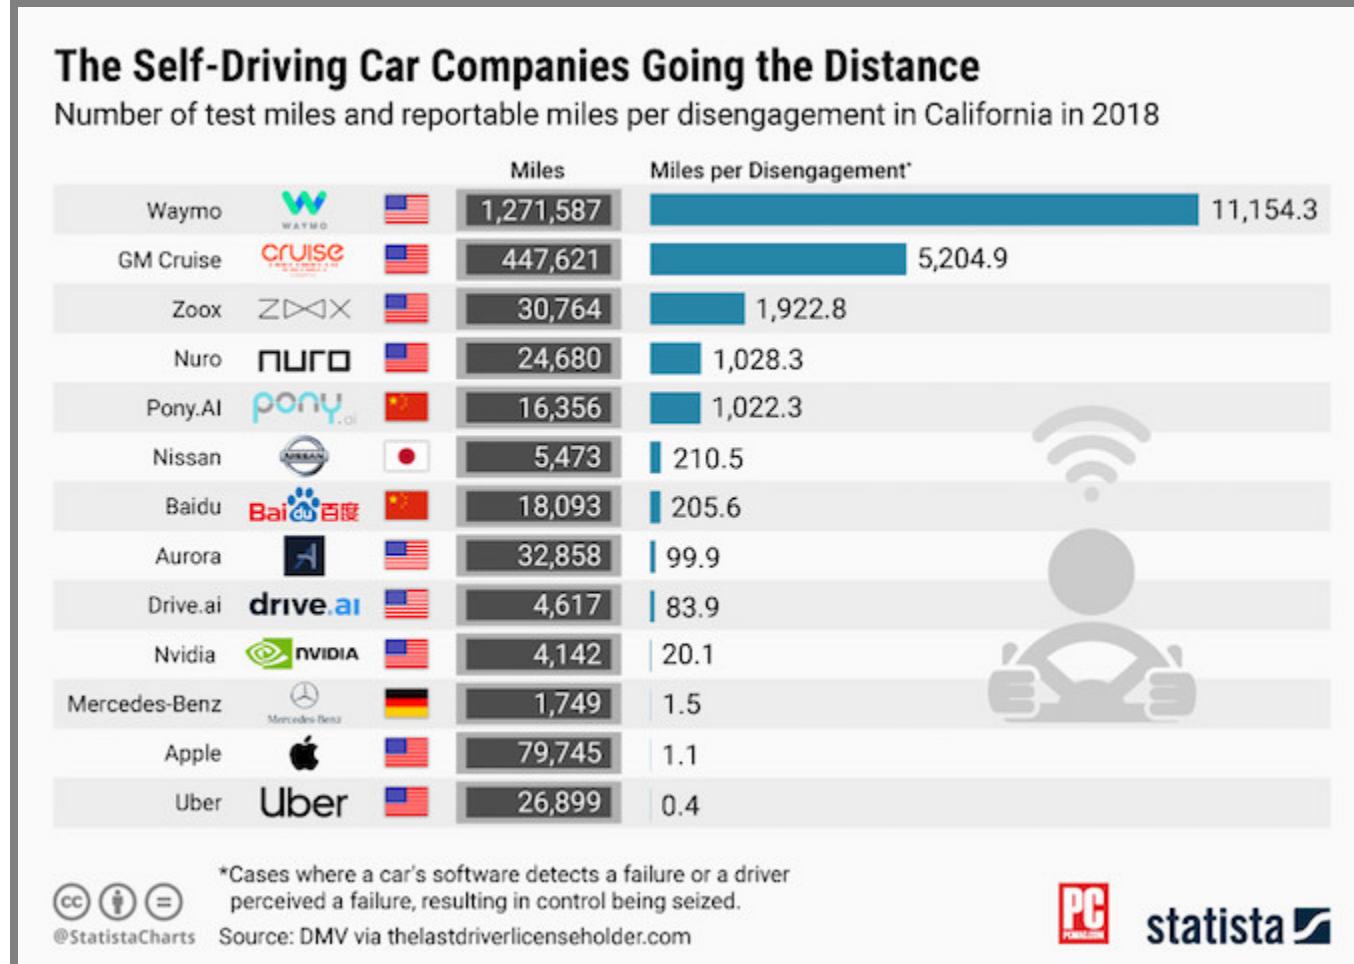
- Extensive field trials
- Usually expensive

**2. Evidence of responsible (safety) engineering process**

- Process with hazard analysis, testing mitigations, etc
- Not sufficient to assure safety

Most standards require both

# Demonstrating Safety

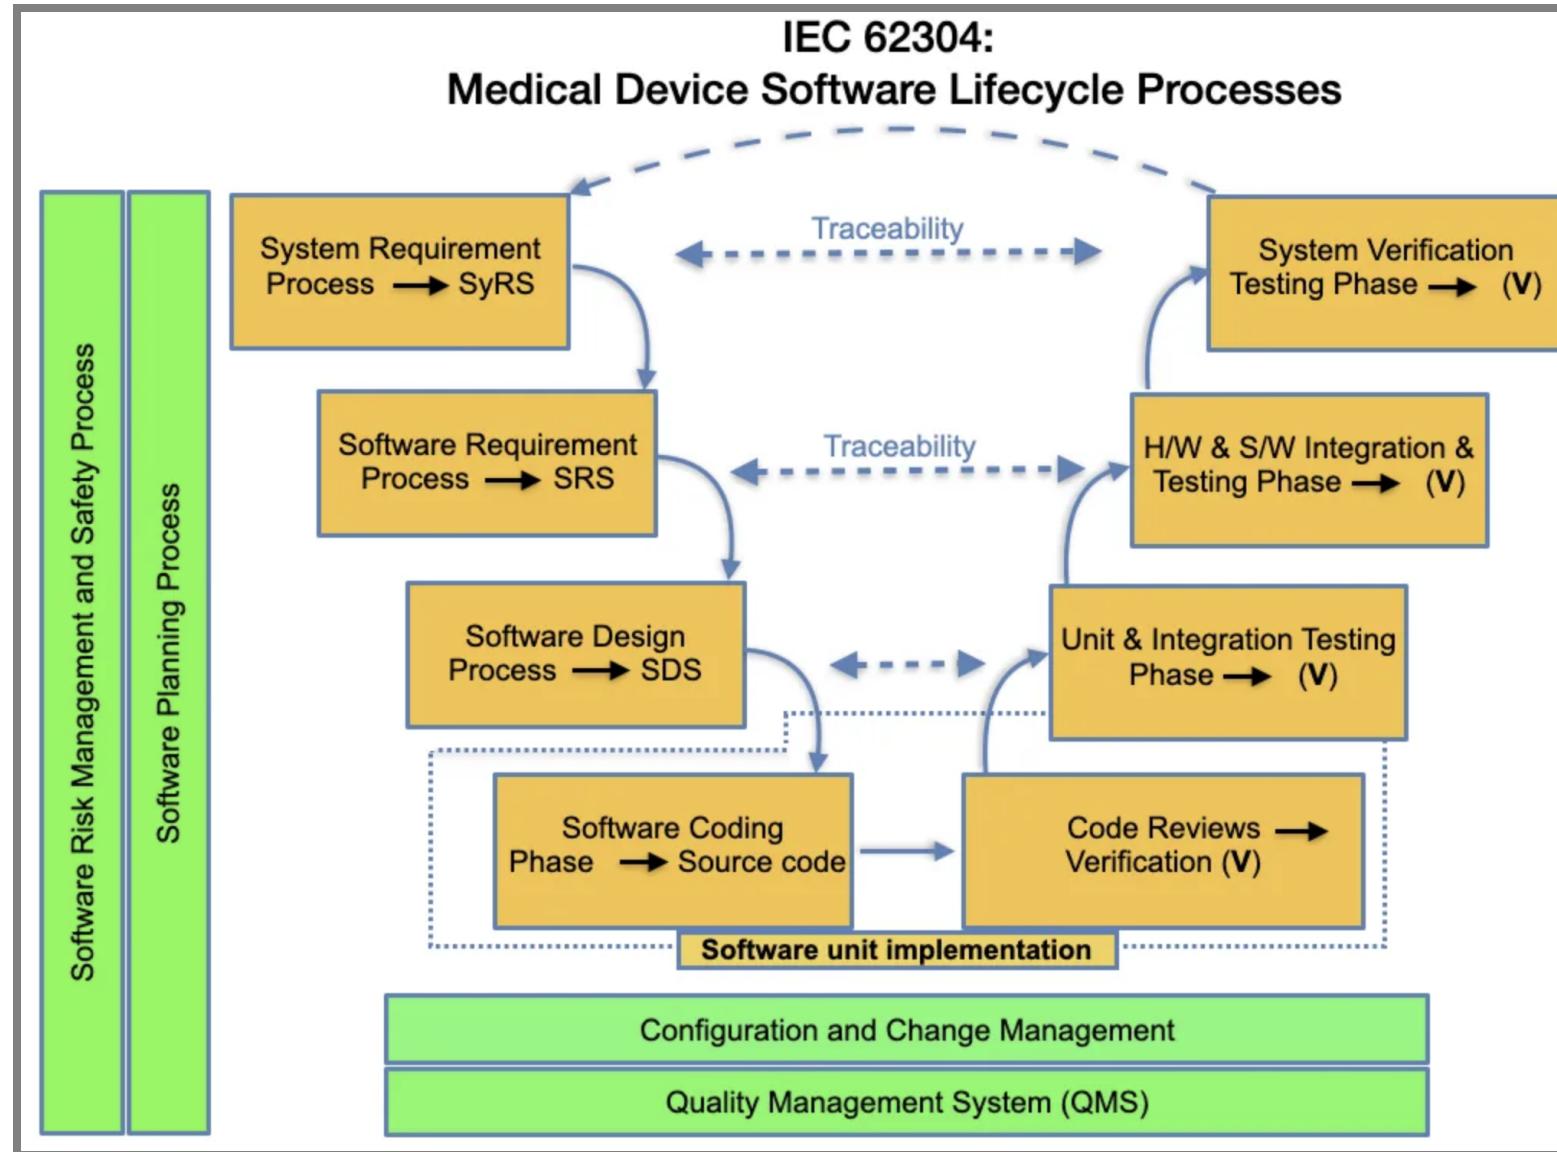


How do we demonstrate to a third-party that our system is safe?

# Safety & Certification Standards

- Guidelines & recommendations for achieving an acceptable level of safety
- Examples: DO-178C (airborne systems), ISO 26262 (automotive), IEC 62304 (medical software), Common Criteria (security)
- Typically, **prescriptive & process-oriented**
  - Recommends use of certain development processes
  - Requirements specification, design, hazard analysis, testing, verification, configuration management, etc.,
- Limitations
  - Most not designed to handle ML systems (exception: UL 4600)
  - Costly to satisfy & certify, but effectiveness unclear (e.g., many FDA-certified products recalled due to safety incidents)
- Good processes are important, but not sufficient; provides only indirect evidence for system safety

# Certification Standards: Example



# Demonstrating Safety

Two main strategies:

1. **Evidence of safe behavior in the field**

- Extensive field trials
- Usually expensive

2. **Evidence of responsible (safety) engineering process**

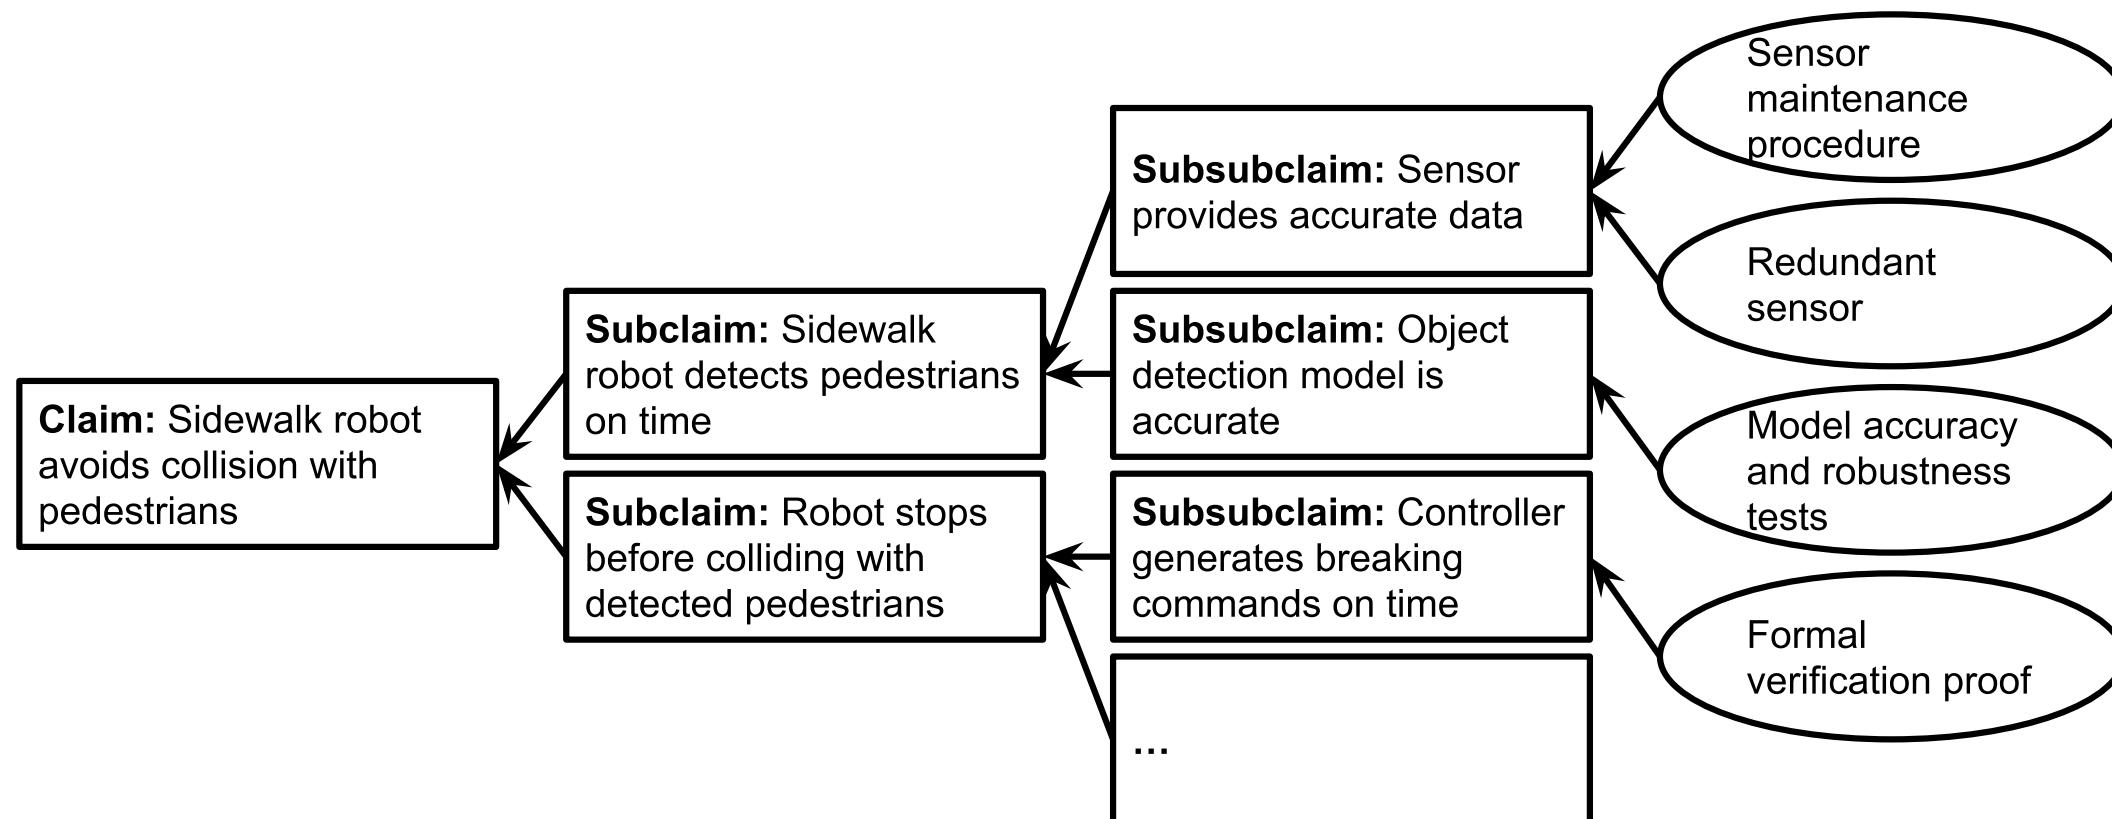
- Process with hazard analysis, testing mitigations, etc
- Not sufficient to assure safety

**Most standards require both, but often not sufficient!**

# Assurance (Safety) Cases

- An emerging approach to demonstrating safety
- An explicit argument that a system achieves a desired safety requirement, along with supporting evidence
- Structure:
  - Argument: A top-level claim decomposed into multiple sub-claims
  - Evidence: Testing, software analysis, formal verification, inspection, expert opinions, design mechanisms...

# Documenting Safety with Assurance (Safety) Cases

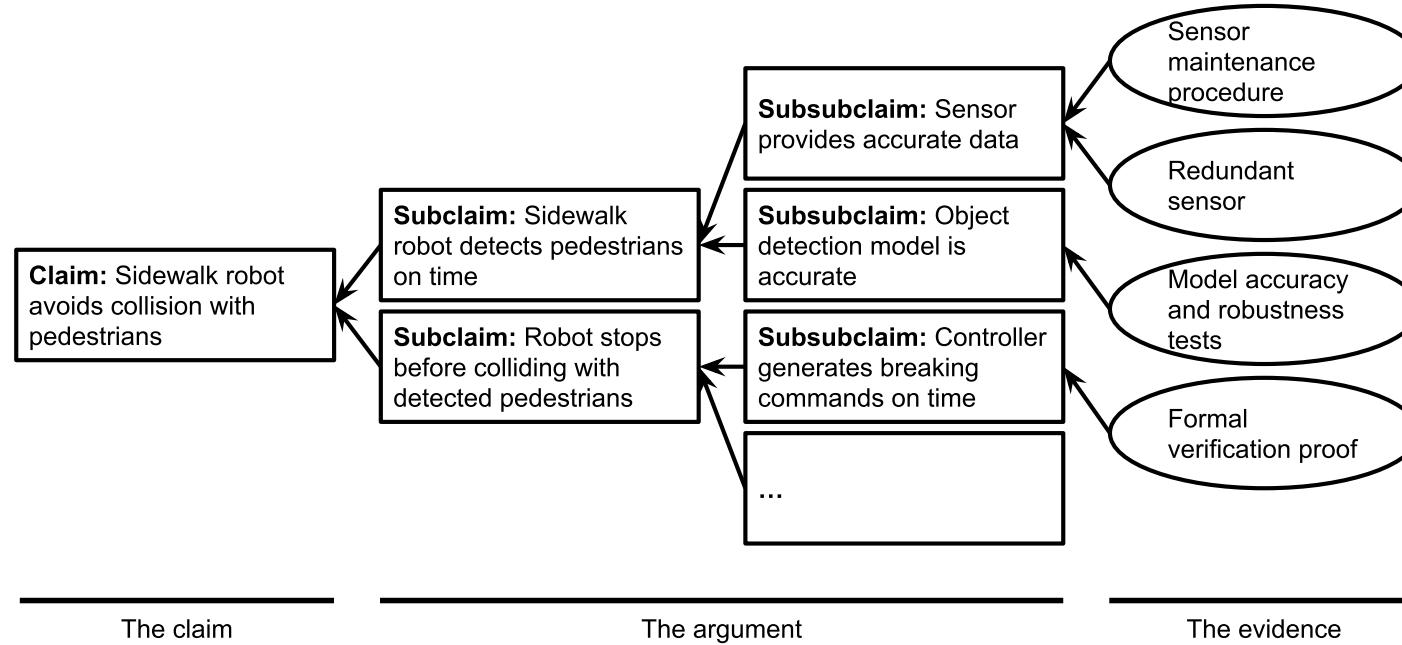


The claim

The argument

The evidence

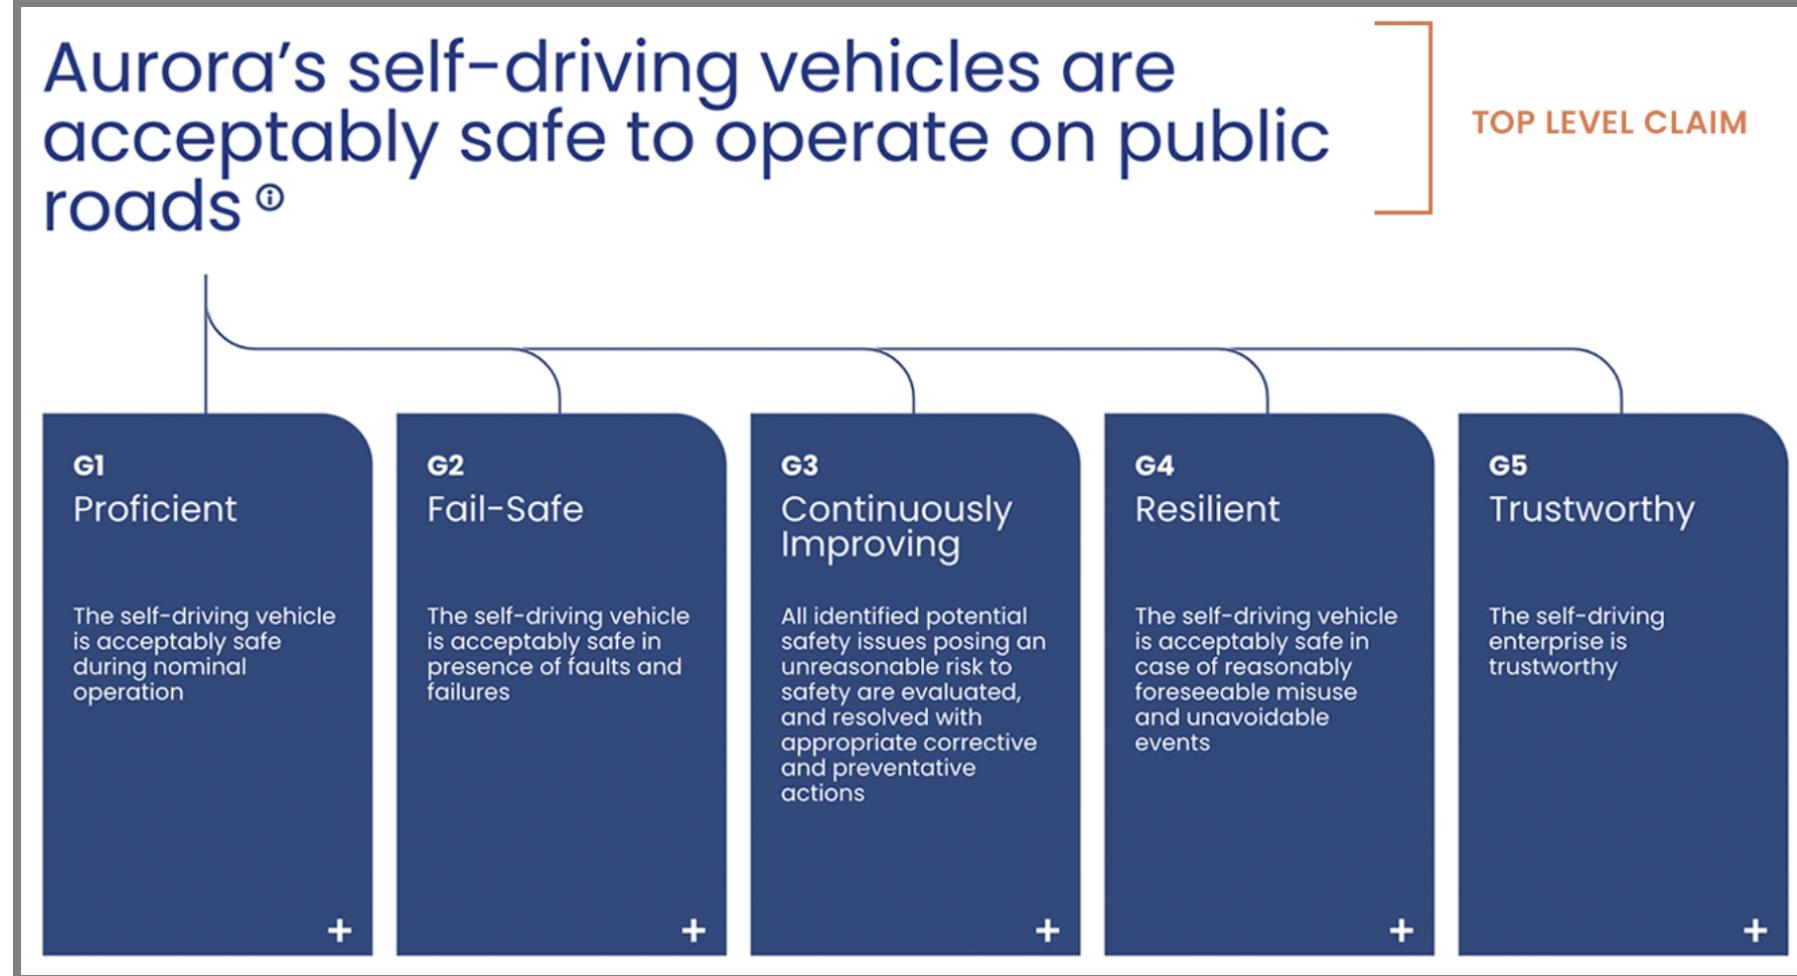
# Assurance Cases: Example



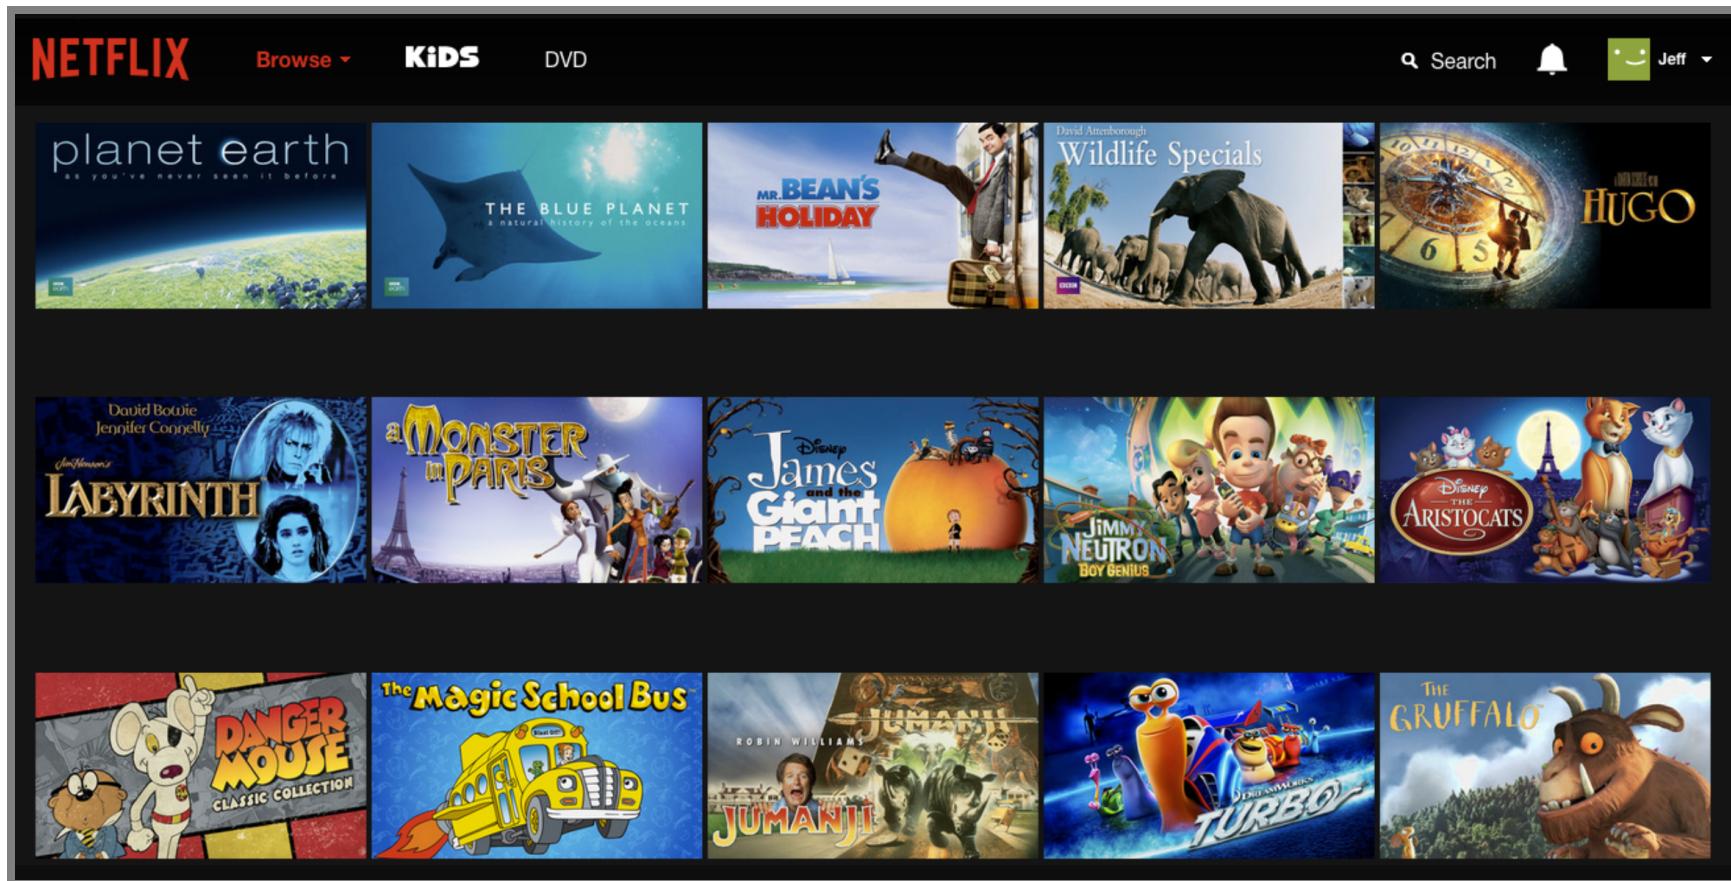
Questions to think about:

- Do sub-claims imply the parent claim?
- Am I missing any sub-claims?
- Is the evidence strong enough to discharge a leaf claim?

# Assurance Cases: Example



# Discussion: Assurance Case for Recommender



# Assurance Cases: Benefits & Limitations

- Provides an explicit structure to the safety argument
  - Easier to navigate, inspect, and refute for third-party auditors
  - Provides traceability between system-level claims & low-level evidence
  - Can also be used for other types of system quality (security, reliability, etc.,)
- Challenges and pitfalls
  - Informal links between claims & evidence, e.g., Does the sub-claims actually imply the top-level claim?
  - Effort in constructing the case & evidence: How much evidence is enough?
  - System evolution: If system changes, must reproduce the case & evidence
- Tools for building & analyzing safety cases available
  - e.g., [ASCE/GSN](#) from Adelard
  - But ultimately, can't replace domain knowledge & critical thinking

# Robustness for ML-based Systems

# Robustness

Environment sometimes **deviates** from expected, normal conditions

- Extreme weathers, unexpected obstacles, etc.,
- Erratic user behaviors; unusually high service demand
- Adversarial actors; users trying to game your system, etc.,

Does your system work reasonably well under these deviations? i.e., is it *robust*?

Most safety-critical systems require some level of robustness

- Not enough to show that system is safe in normal conditions

# Defining Robustness for ML:

- A prediction for input  $x$  is robust if the outcome is stable under minor perturbations to the input:
  - $\forall x'. d(x, x') < \epsilon \Rightarrow f(x) = f(x')$
  - distance function  $d$  and permissible distance  $\epsilon$  depends on the problem domain!
- A model is said to be robust if most predictions are robust
- An important concept in safety and security settings
  - In safety, perturbations tend to be random or predictable (e.g., sensor noise due to weather conditions)
  - In security, perturbations are intentionally crafted (e.g., adversarial attacks)

# Robustness and Distance for Images

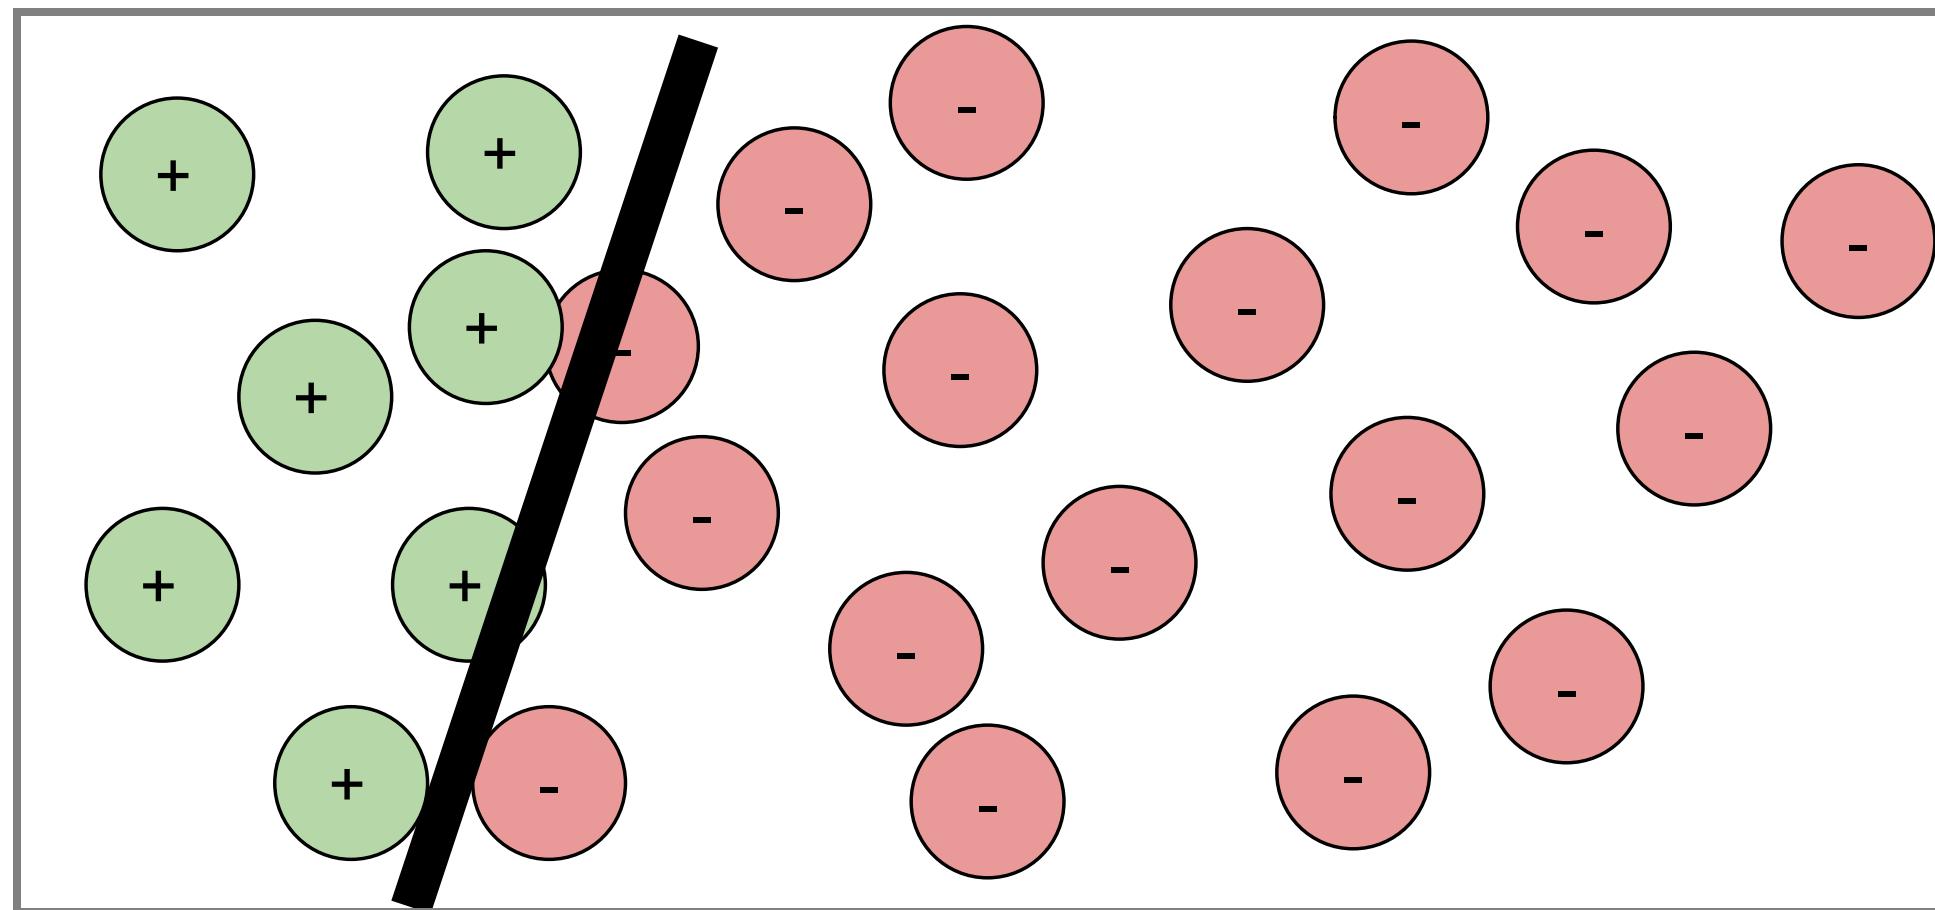
- Slight rotation, stretching, or other transformations
- Change many pixels minimally (below human perception)
- Change only few pixels
- Change most pixels mostly uniformly, e.g., brightness

Attack	Original	Lower	Upper
$L_\infty$			
Rotation			

≡ Image: [An abstract domain for certifying neural networks](#). Gagandeep et al., POPL (2019).

# No Model is Fully Robust

- Every useful model has at least one decision boundary
- Predictions near that boundary are not (and should not) be robust



# Robustness of Interpretable Models

Is this model robust?

Is the prediction for a 20 year old male with 2 priors robust?

```
IF age between 18-20 and sex is male THEN predict arrest  
ELSE  
IF age between 21-23 and 2-3 prior offenses THEN predict arres  
ELSE  
IF more than three priors THEN predict arrest  
ELSE predict no arrest
```

# Evaluating ML Robustness

- Lots of on-going research (especially for DNNs)
- Formal verification
  - Constraint solving or abstract interpretation over computations in neuron activations
  - Conservative abstraction, may label robust inputs as not robust
  - Currently not very scalable
  - Example: *An abstract domain for certifying neural networks*. Gagandeep et al., POPL (2019).
- Sampling
  - Sample within distance, compare prediction to majority prediction
  - Probabilistic guarantees possible (with many queries, e.g., 100k)
  - Example: *Certified adversarial robustness via randomized smoothing*. Cohen, Rosenfeld, and Kolter, ICML (2019).

# ML Robustness: Limitations

- Lots of on-going research (especially for DNNs)
- Mostly input-centric, focusing on small ( $\epsilon$ ) perturbations
  - Common use case: Robustness against adversarial attacks
  - Q. But do these perturbations matter for safety?
- In practice: Perturbations result from environmental changes!
  - Which parts of the world does the software sense?
  - Can those parts change over time? Can the sensors be noisy, faulty, etc.,? (these are **domain-specific**)
  - What input perturbations could be caused by from these changes/noise...?

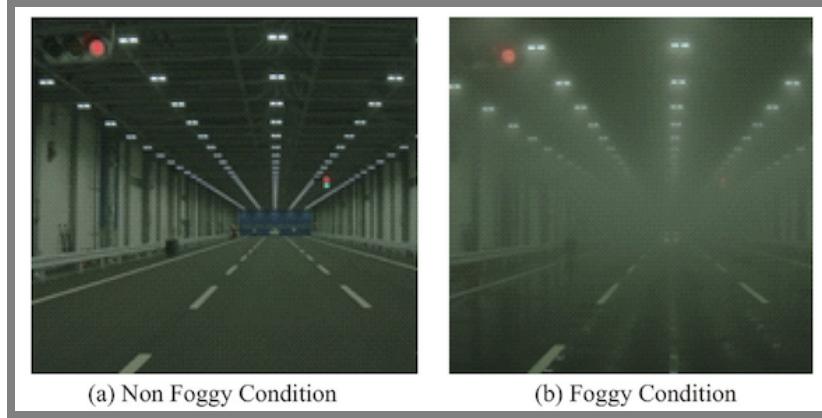
# Robustness in a Safety Setting

- Does the model detect stop signs under normal settings?
- Does the model detect stop signs under deviations?
  - Poor lighting? In fog? With a tilted camera? Sensor noise?
  - With stickers taped to the sign? (Does it matter?)



Image: David Silver. [Adversarial Traffic Signs](#). Blog post, 2017

# Improving Robustness for Safety

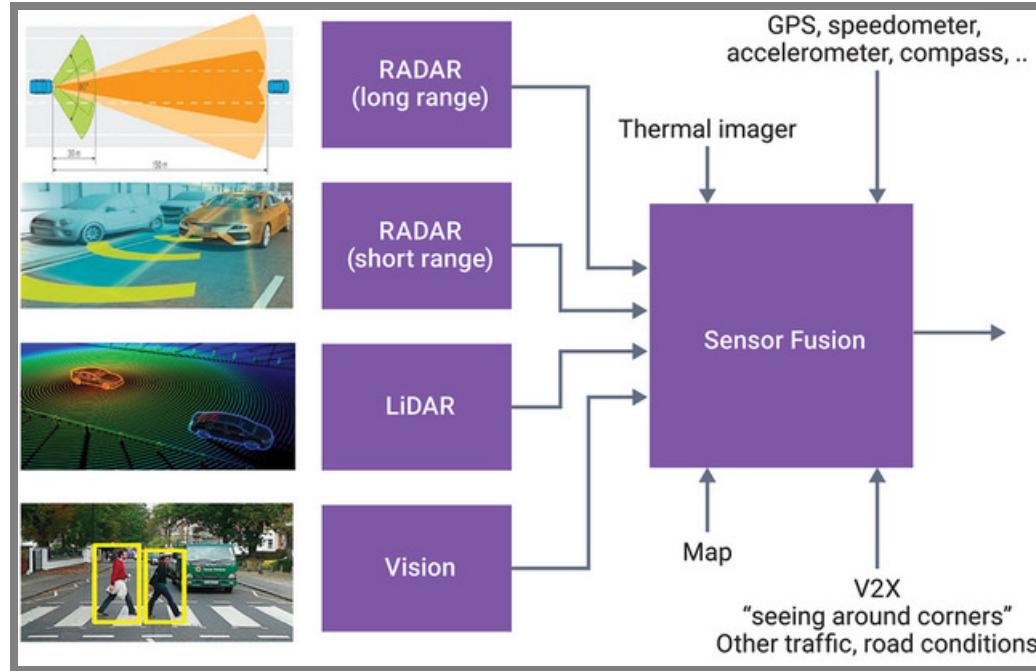


Learn more robust models

- Test/think about domain-specific deviations that might result in perturbations to model input (e.g., fogs, snow, sensor noise)
- Curate data for those abnormal scenarios or augment training data with transformed inputs

 *Automated driving recognition technologies for adverse weather conditions.* Yoneda et al., (2019).

# Improving Robustness for Safety



## Design mechanisms

- Deploy redundant components for critical tasks (e.g., vision + map)
- Ensemble learning: Combine models with different biases
- Multiple, independent sensors (e.g., LiDAR + radar + cameras)

# Improving Robustness for Safety

## Design mechanisms

- Deploy redundant components for critical tasks (e.g., vision + map)
- Ensemble learning: Combine models with different biases
- Multiple, independent sensors (e.g., LiDAR + radar + cameras)

## Robustness checking at inference time

- Handle inputs with non-robust predictions differently (e.g. discard or output low confidence score for outliers)
- Downside: Raises cost of prediction; may not be suitable for time-sensitive applications (e.g., self-driving cars)

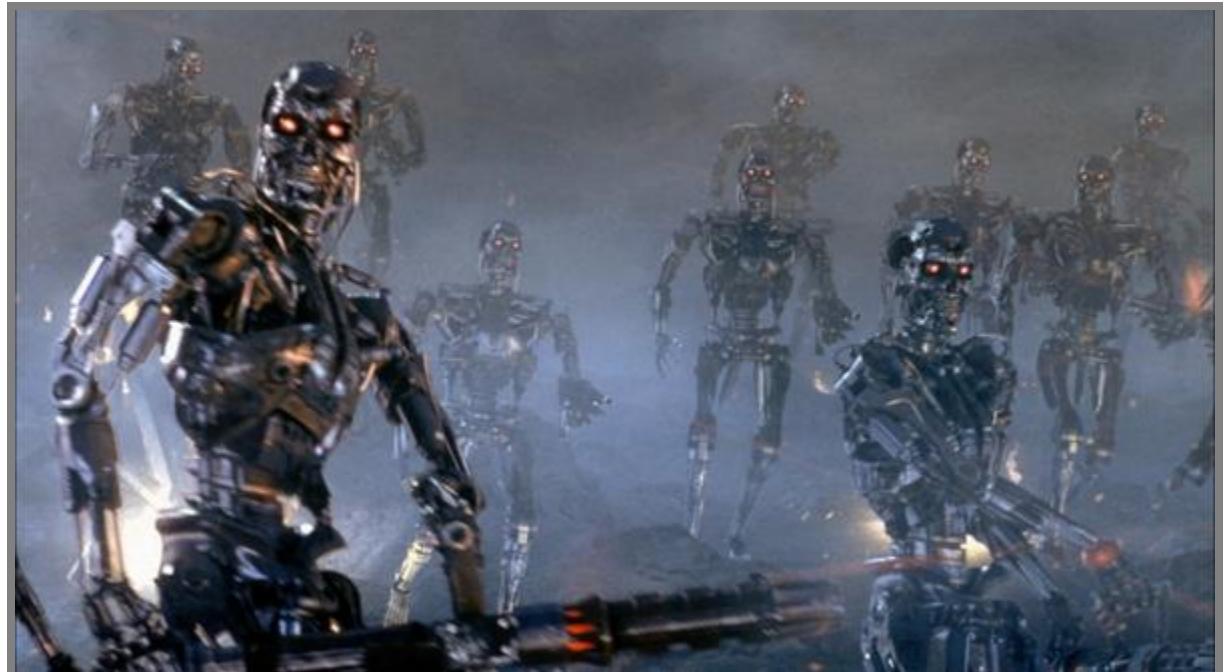
# Breakout: Robustness

Scenario: Medical use of transcription service, dictate diagnoses and prescriptions

As a group, tagging members, post to #lecture:

1. *What safety concerns can you anticipate?*
2. *What deviations are you concerned about?*
3. *How would improve the robustness of the overall system?*

# AI Safety



Amodei, Dario, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané.  
≡ "Concrete problems in AI safety." arXiv preprint arXiv:1606.06565 (2016).

# Your Favorite AI Dystopia?



# The AI Alignment Problem

AI is optimized for a specific objective/cost function

- Inadvertently cause undesirable effects on the environment
- e.g., **Transport robot**: Move a box to a specific destination
- Side effects: Scratch furniture, bump into humans, etc.,

Side effects may cause ethical/safety issues (e.g., social media optimizing for clicks, causing teen depression)

Difficult to define sensible fitness functions:

- Perform X *subject to common-sense constr. on the environment*
- Perform X *but avoid side effects to the extent possible*

# Reward Hacking

*PlayFun algorithm pauses the game of Tetris indefinitely to avoid losing*

*When about to lose a hockey game, the PlayFun algorithm exploits a bug to make one of the players on the opposing team disappear from the map, thus forcing a draw.*

*Self-driving car rewarded for speed learns to spin in circles*

Example: Coast Runner

# Reward Hacking

- AI can be good at finding loopholes to achieve a goal in unintended ways
- Technically correct, but does not follow *designer's informal intent*
- Many possible causes, incl. partially observed goals, abstract rewards, feedback loops
- In general, a very challenging problem!
  - Difficult to specify goal & reward function to avoid all possible hacks
  - Requires careful engineering and iterative reward design

# Reward Hacking -- Many Examples



Victoria Krakovna  
@vkrakovna · [Follow](#)



New resource: a master list of examples of AI systems gaming their objective specification:  
[docs.google.com/spreadsheets/d...](https://docs.google.com/spreadsheets/d/)

Accompanying blog post:  
[vkrakovna.wordpress.com/2018/04/02/spe...](http://vkrakovna.wordpress.com/2018/04/02/spe...)

Thanks [@gwern](#) and [@catherineols](#) for the inspiration and feedback on putting this together!



[vkrakovna.wordpress.com](http://vkrakovna.wordpress.com)  
Specification gaming examples in AI  
Update: for a more detailed introduction to specification ga...

12:37 PM · Apr 2, 2018



[Read the full conversation on Twitter](#)

# Exploiting Human Weakness

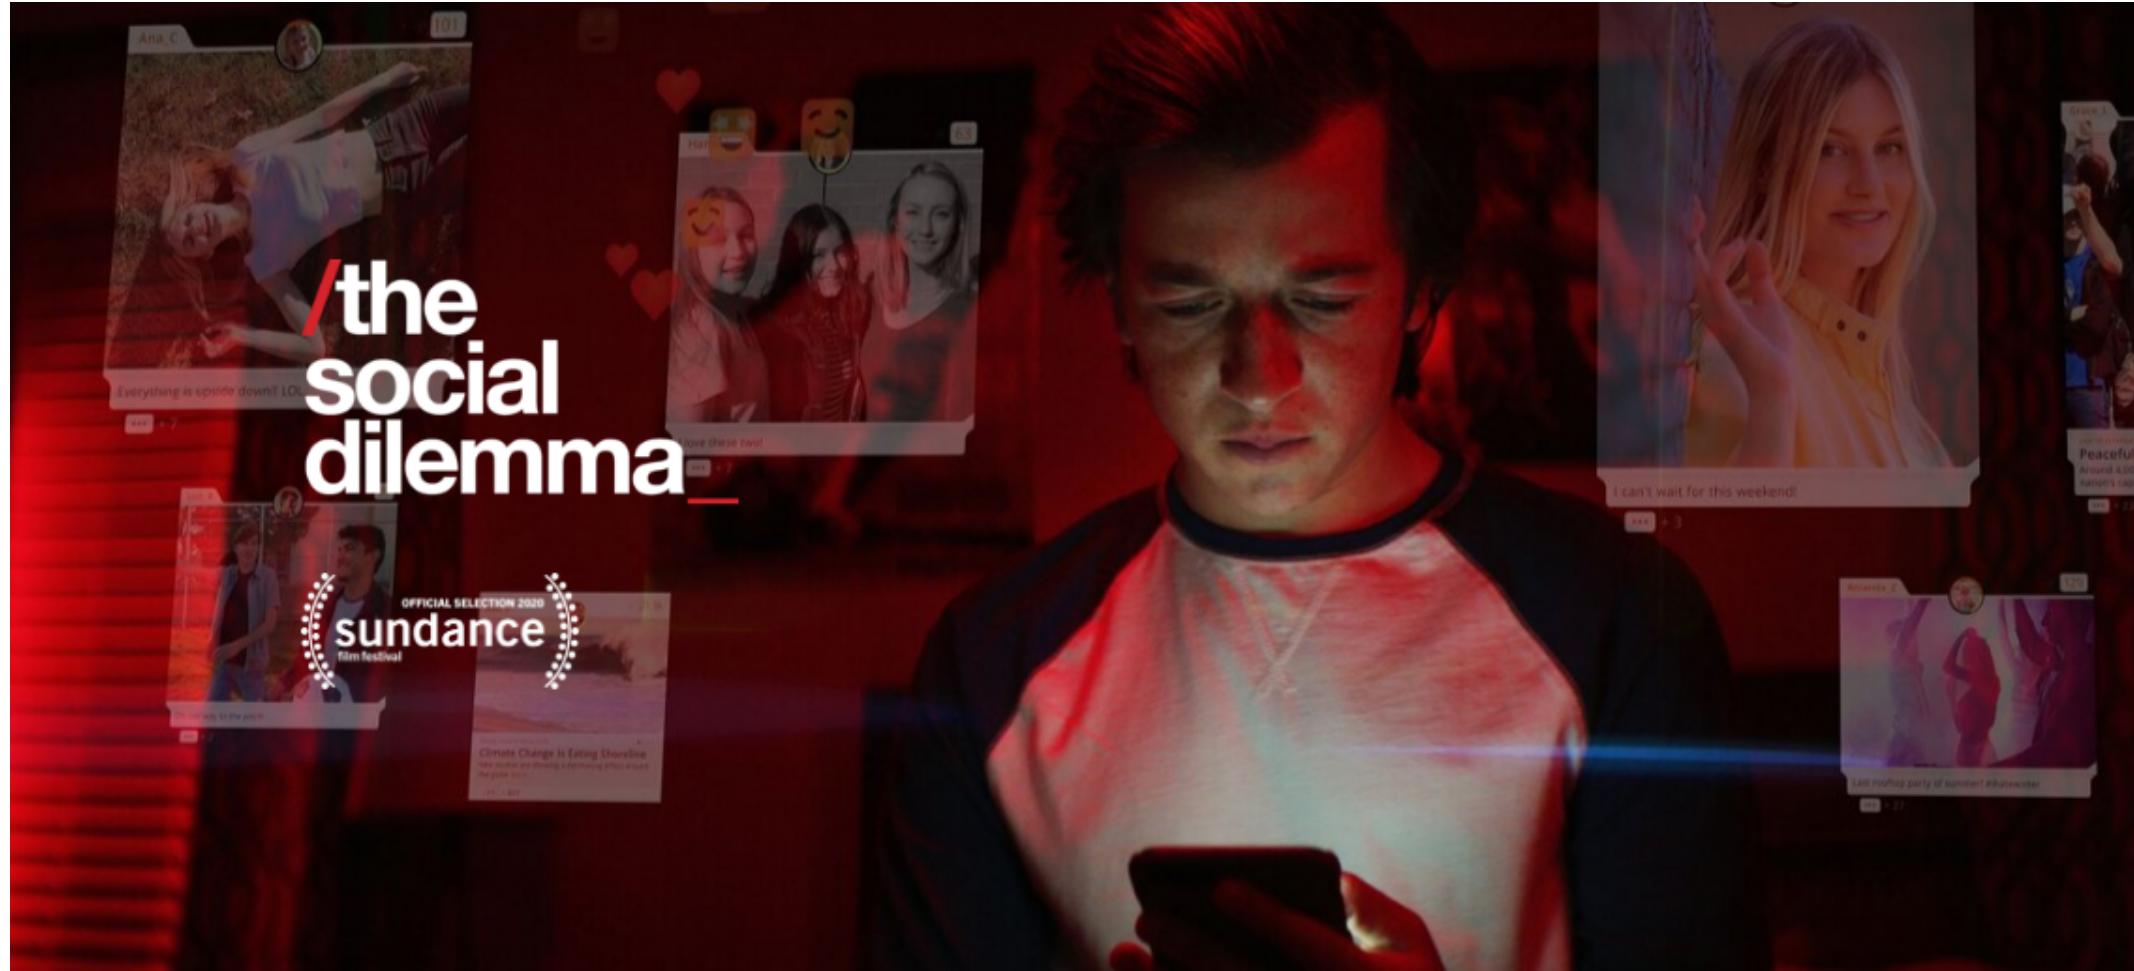
The Morality Of A/B Testing

Josh Constine @joshconstine 11:50 PM EDT • June 29, 2014



We don't use the "real" Facebook. Or Twitter. Or Google, Yahoo,

# Exploiting Human Weakness



≡ See also Center for Humane Technology

# AI Alignment Problem = Requirements Problem

Recall: "World vs. machine"

- Identify stakeholders in the environment & possible effects on them
- Anticipate side effects, feedback loops
- Constrain the scope of the system
- Perfect contracts usually infeasible, undesirable

But more requirements engineering unlikely to be only solution

# Other Challenges

- Safe Exploration
  - Exploratory actions "in production" may have consequences
  - e.g., trap robots, crash drones
- Robustness to Drift
  - Drift may lead to poor performance that may not even be recognized
- Scalable Oversight
  - Cannot provide human oversight over every action (or label all possible training data)
  - Use indirect proxies in telemetry to assess success/satisfaction

# Existential AI Risk

Existential risk and AI alignment common in research

Funding through *longtermism* branch of effective altruism

*(Longtermism is the view that positively influencing the longterm future is a key moral priority of our time.)*

Ord estimates 10% existential risk from unaligned AI in 100 years

**Our view:** AI alignment not a real concern for the kind of ML-enabled products we consider here

## Speaker notes

Relevant for reinforcement learning and AGI



# More pressing AI risks?

AI

## Ethicists fire back at 'AI Pause' letter they say 'ignores the actual harms'

Devin Coldewey @techcrunch / 8:18 PM EDT • March 31, 2023

*"Those hypothetical risks are the focus of a dangerous ideology called longtermism that ignores the actual harms resulting from the deployment of AI systems today," they wrote, citing worker exploitation, data theft, synthetic media that props up existing power structures and the further concentration of those power structures in fewer hands.*

# Practical Alignment Problems

Does the model goal align with the system goal? Does the system goal align with the user's goals?

- Profits (max. accuracy) vs fairness
- Engagement (ad sales) vs enjoyment, mental health
- Accuracy vs operating costs

Test model *and* system quality *in production*

(see requirements engineering and architecture lectures)

# Beyond Traditional Safety Critical Systems

# Beyond Traditional Safety Critical Systems

- Recall: Legal vs ethical
- Safety analysis not only for regulated domains (nuclear power plants, medical devices, planes, cars, ...)
- Many end-user applications have a safety component

Q. Examples?

# Mental Health

≡ 🔎

**healthline**

SUBSCRIBE

HEALTH NEWS ✓ Fact Checked

## The FOMO Is Real: How Social Media Increases Depression and Loneliness

Written by [Gigen Mammoser](#) on December 10, 2018

New research reveals how social media platforms like Facebook can greatly affect your mental health.



≡

# Mental Health

≡ 🔎

**healthline**

SUBSCRIBE

HEALTH NEWS ✓ Fact Checked

## The FOMO Is Real: How Social Media Increases Depression and Loneliness

Written by [Gigen Mammoser](#) on December 10, 2018

New research reveals how social media platforms like Facebook can greatly affect your mental health.



≡

# IoT

The @netatmo servers are down and twitter is already full of freezing people not able to control their heating :D (via [protected]) / cc @internetofshit

eran  
DivemasterK

no Are your servers down? I can't get my app to turn on heating at 1:02 from Wicklow, Ireland

Kiran vadgama @kiran\_vadgama  
netatmo hi my manual override option on the smart thermostat is not working and when i open the app it comes up with an error message saying the servers are down. Can i override a manual setting?  
1.18, 20:58

Andy Mc @ITakeSugar  
Replying to @leviseedaniel and @tyrestighe Is there a way to control the heating when the servers are down, it's freezing right now.  
22.11.18, 20:38

james brun @jamesbrun · 1 reply · 1 retweet · 1 like  
Replies to @tyrestighe @leviseedaniel  
netatmo issue. Can't control heating via app or website. Can't login to netatmo.com or control from there. What is going on with netatmo ?

3:15 PM - 22 Nov 2018

1,659 Retweets 2,280 Likes

≡

# Addiction

NO MERCY NO MALICE

## Robinhood Has Gamified Online Trading Into an Addiction

Tech's obsession with addiction will hurt us all



Scott Galloway [Follow](#)

Jun 23 · 7 min read ★



*Warning: This post contains a discussion of suicide.*

**A**ddiction is the inability to stop consuming a chemical or pursuing an activity although it's causing harm.

I engage with almost every substance or behavior associated with addiction: alcohol, drugs, coffee, porn, sex, gambling, work, spending,

# Society: Unemployment Engineering / Deskilling



## Speaker notes

The dangers and risks of automating jobs.

Discuss issues around automated truck driving and the role of jobs.

See for example: Andrew Yang. The War on Normal People. 2019

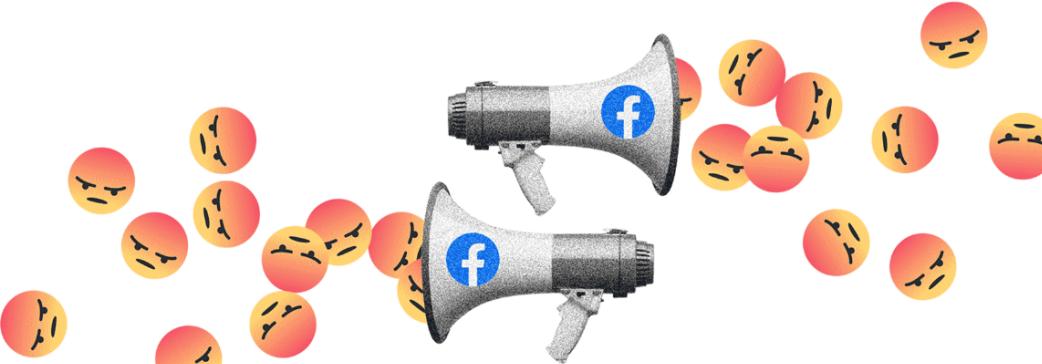


# Society: Polarization

THE WALL STREET JOURNAL.

SUBSCRIBE

SIGN IN



TECH

## Facebook Executives Shut Down Efforts to Make the Site Less Divisive

The social-media giant internally studied how it polarizes users, then largely shelved the research

By [Jeff Horwitz](#) and [Deepa Seetharaman](#)

May 26, 2020 11:38 am ET

## Speaker notes

Recommendations for further readings: <https://www.nytimes.com/column/kara-swisher>, <https://podcasts.apple.com/us/podcast/recode-decode/id1011668648>

Also isolation, Cambridge Analytica, collaboration with ICE, ...



# Environmental: Energy Consumption

[≡](#)

**NewScientist**

[SUBSCRIBE AND SAVE 69%](#)

[≡](#)

## Creating an AI can be five times worse for the planet than a car

[!\[\]\(a08936a476c8059a8f420096a8be987e\_img.jpg\)](#) [!\[\]\(7bd4a587ec06429114bbc976b29a75db\_img.jpg\)](#) [!\[\]\(5a0e8369b88ff75d66ff4f43b2594d65\_img.jpg\)](#) [!\[\]\(04020706e9c327c9c4f9c6dbe28b2a57\_img.jpg\)](#) [!\[\]\(af4a9bd1dd4570b080694398609ea175\_img.jpg\)](#) [!\[\]\(75762ac3cc680b204b6418b66da1fa60\_img.jpg\)](#) [!\[\]\(cfba3c03054f1ab5fa06ba90731670fb\_img.jpg\)](#)

TECHNOLOGY 6 June 2019

By [Donna Lu](#)



64

# Exercise

*Look at apps on your phone. Which apps have a safety risk and use machine learning?*

Consider safety broadly: including stress, mental health, discrimination, and environment pollution

# Takeaway

- Many systems have safety concerns
- ... not just nuclear power plants, planes, cars, and medical devices
- Do the right thing, even without regulation
- Consider safety broadly: including stress, mental health, discrimination, and environment pollution
- Start with requirements and hazard analysis

# Designing for Safety

See Lecture Planning for Mistakes

# Safety Assurance with ML Components

- Consider ML components as unreliable, at most probabilistic guarantees
- Testing, testing, testing (+ simulation)
  - Focus on data quality & robustness
- *Adopt a system-level perspective!*
- Consider safe system design with unreliable components
  - Traditional systems and safety engineering
  - Assurance cases
- Understand the problem and the hazards
  - System level, goals, hazard analysis, world vs machine
  - Specify *end-to-end system behavior* if feasible

# Summary

- Defining safety: absence of harm to people, property, and environment -- consider broadly; safety  $\neq$  reliability
- *Adopt a safety mindset!*
- Assume all components will eventually fail in one way or another, especially ML components
- Hazard analysis to identify safety risks and requirements; classic safety design at the system level
- Robustness: Identify & address relevant deviations
- AI alignment: AI goals are difficult to specify precisely; susceptible to negative side effect & reward hacking

# Further Readings

- Borg, Markus, Cristofer Englund, Krzysztof Wnuk, Boris Duran, Christoffer Levandowski, Shenjian Gao, Yanwen Tan, Henrik Kaijser, Henrik Lönn, and Jonas Törnqvist. “[Safely entering the deep: A review of verification and validation for machine learning and a challenge elicitation in the automotive industry.](#)” Journal of Automotive Software Engineering. 2019
- Leveson, Nancy G. [Engineering a safer world: Systems thinking applied to safety.](#) The MIT Press, 2016.
- Salay, Rick, and Krzysztof Czarnecki. “[Using machine learning safely in automotive software: An assessment and adaption of software process requirements in ISO 26262.](#)” arXiv preprint arXiv:1808.01614 (2018).
- Mohseni, Sina, Mandar Pitale, Vasu Singh, and Zhangyang Wang. “[Practical Solutions for Machine Learning Safety in Autonomous Vehicles.](#)” SafeAI workshop at AAAI’20, (2020).
- Huang, Xiaowei, Daniel Kroening, Wenjie Ruan, James Sharp, Youcheng Sun, Emese Thamo, Min Wu, and Xinping Yi. “[A survey of safety and trustworthiness of deep neural networks: Verification, testing, adversarial attack and defence, and interpretability.](#)” Computer Science Review 37 (2020).
- Amodei, Dario, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. “[Concrete problems in AI safety.](#)” arXiv preprint arXiv:1606.06565 (2016).

